



7 characteristics of container-native infrastructure



Joyent®

container-native all the things

**HI!
MY NAME IS**

Casey Bisson

Public Cloud

Triton Elastic Container Service. We run our customer's mission critical applications on container native infrastructure

Private Cloud

Triton Elastic Container Infrastructure is an on-premise, container run-time environment used by some of the world's most recognizable brands



Public Cloud

Triton Elastic

customer's microservices
container native



joyent / sdc

Main SmartDataCenter Project — Edit

97 commits · 2 branches · 0 releases · 17 contributors

branch: master → sdc / +

RELENG-613: deprecate sdc-zookeeper
trentm authored 5 days ago

assets · arch diagram

RELENG-613: deprecate sdc-zookeeper
use the same form for the git urls as the other repos
setup.bat: VMware Workstation path.
update repos.md, start reviewing notes

TOOLS-607 I for one, welcome my new open source license
joyent/sdc#33: make clone-all-repos
doc: fix download link to coal~setup.bat .

TOOLS-623 Add repos list

latest commit 40c7667838 · 6 months ago

5 days ago · 6 months ago

25 days ago · 8 months ago

25 days ago · 8 months ago

5 months ago · 25 days ago

8 months ago · 8 months ago

5 months ago · 25 days ago

8 months ago · 8 months ago

5 months ago · 25 days ago

8 months ago · 8 months ago

it's open source!

fork me, pull me: <https://github.com/joyent/sdc>

optimized to deliver next

Unwatch 83 · Unstar 346 · Fork

Code · Issues · Pull requests · Wiki · Pulse · Graphs · Settings · SSH clone URL · Clone · Download



The best place to run Docker



Portability

From laptop to any
public or private cloud

Great for DevOps

Tools for management,
deployment & scale

Productivity

Faster code, test
and deploy

some
background

Infrastructure containers

- We use them like VMs
- They perform like bare metal
- We can pack 3x more containers than VMs on a single compute node
- They require no changes in application architecture
- First appeared in FreeBSD as jails in 1999-ish; defined for Solaris in 2002; Linux implementation in progress (user namespaces mainlined in 2013)
- Joyent has been using them to power our public cloud for almost a decade

Virtualization and Namespace Isolation in the Solaris Operating System (PSARC/2002/174)

John Beck, David Comay, Ozgur L., Daniel Price, and Andy T.
Solaris Kernel Technology

Andrew G. and Blaise S.
Solaris Network and Security Technologies

Revision 1.6 (OpenSolaris)

September 7, 2006

SmartOS container hypervisor

- Forked from Open Solaris in 2010
- Part of the Illumos community
- Open source
- Includes the core compute, network, and storage virtualization components:
 - **Zones** = truly secure containers
 - **Crossbow** = in-kernel network virtualization, offers unique NIC(s) for each container
 - **ZFS** = fast and secure virtualized filesystems

Triton (was SmartDataCenter)

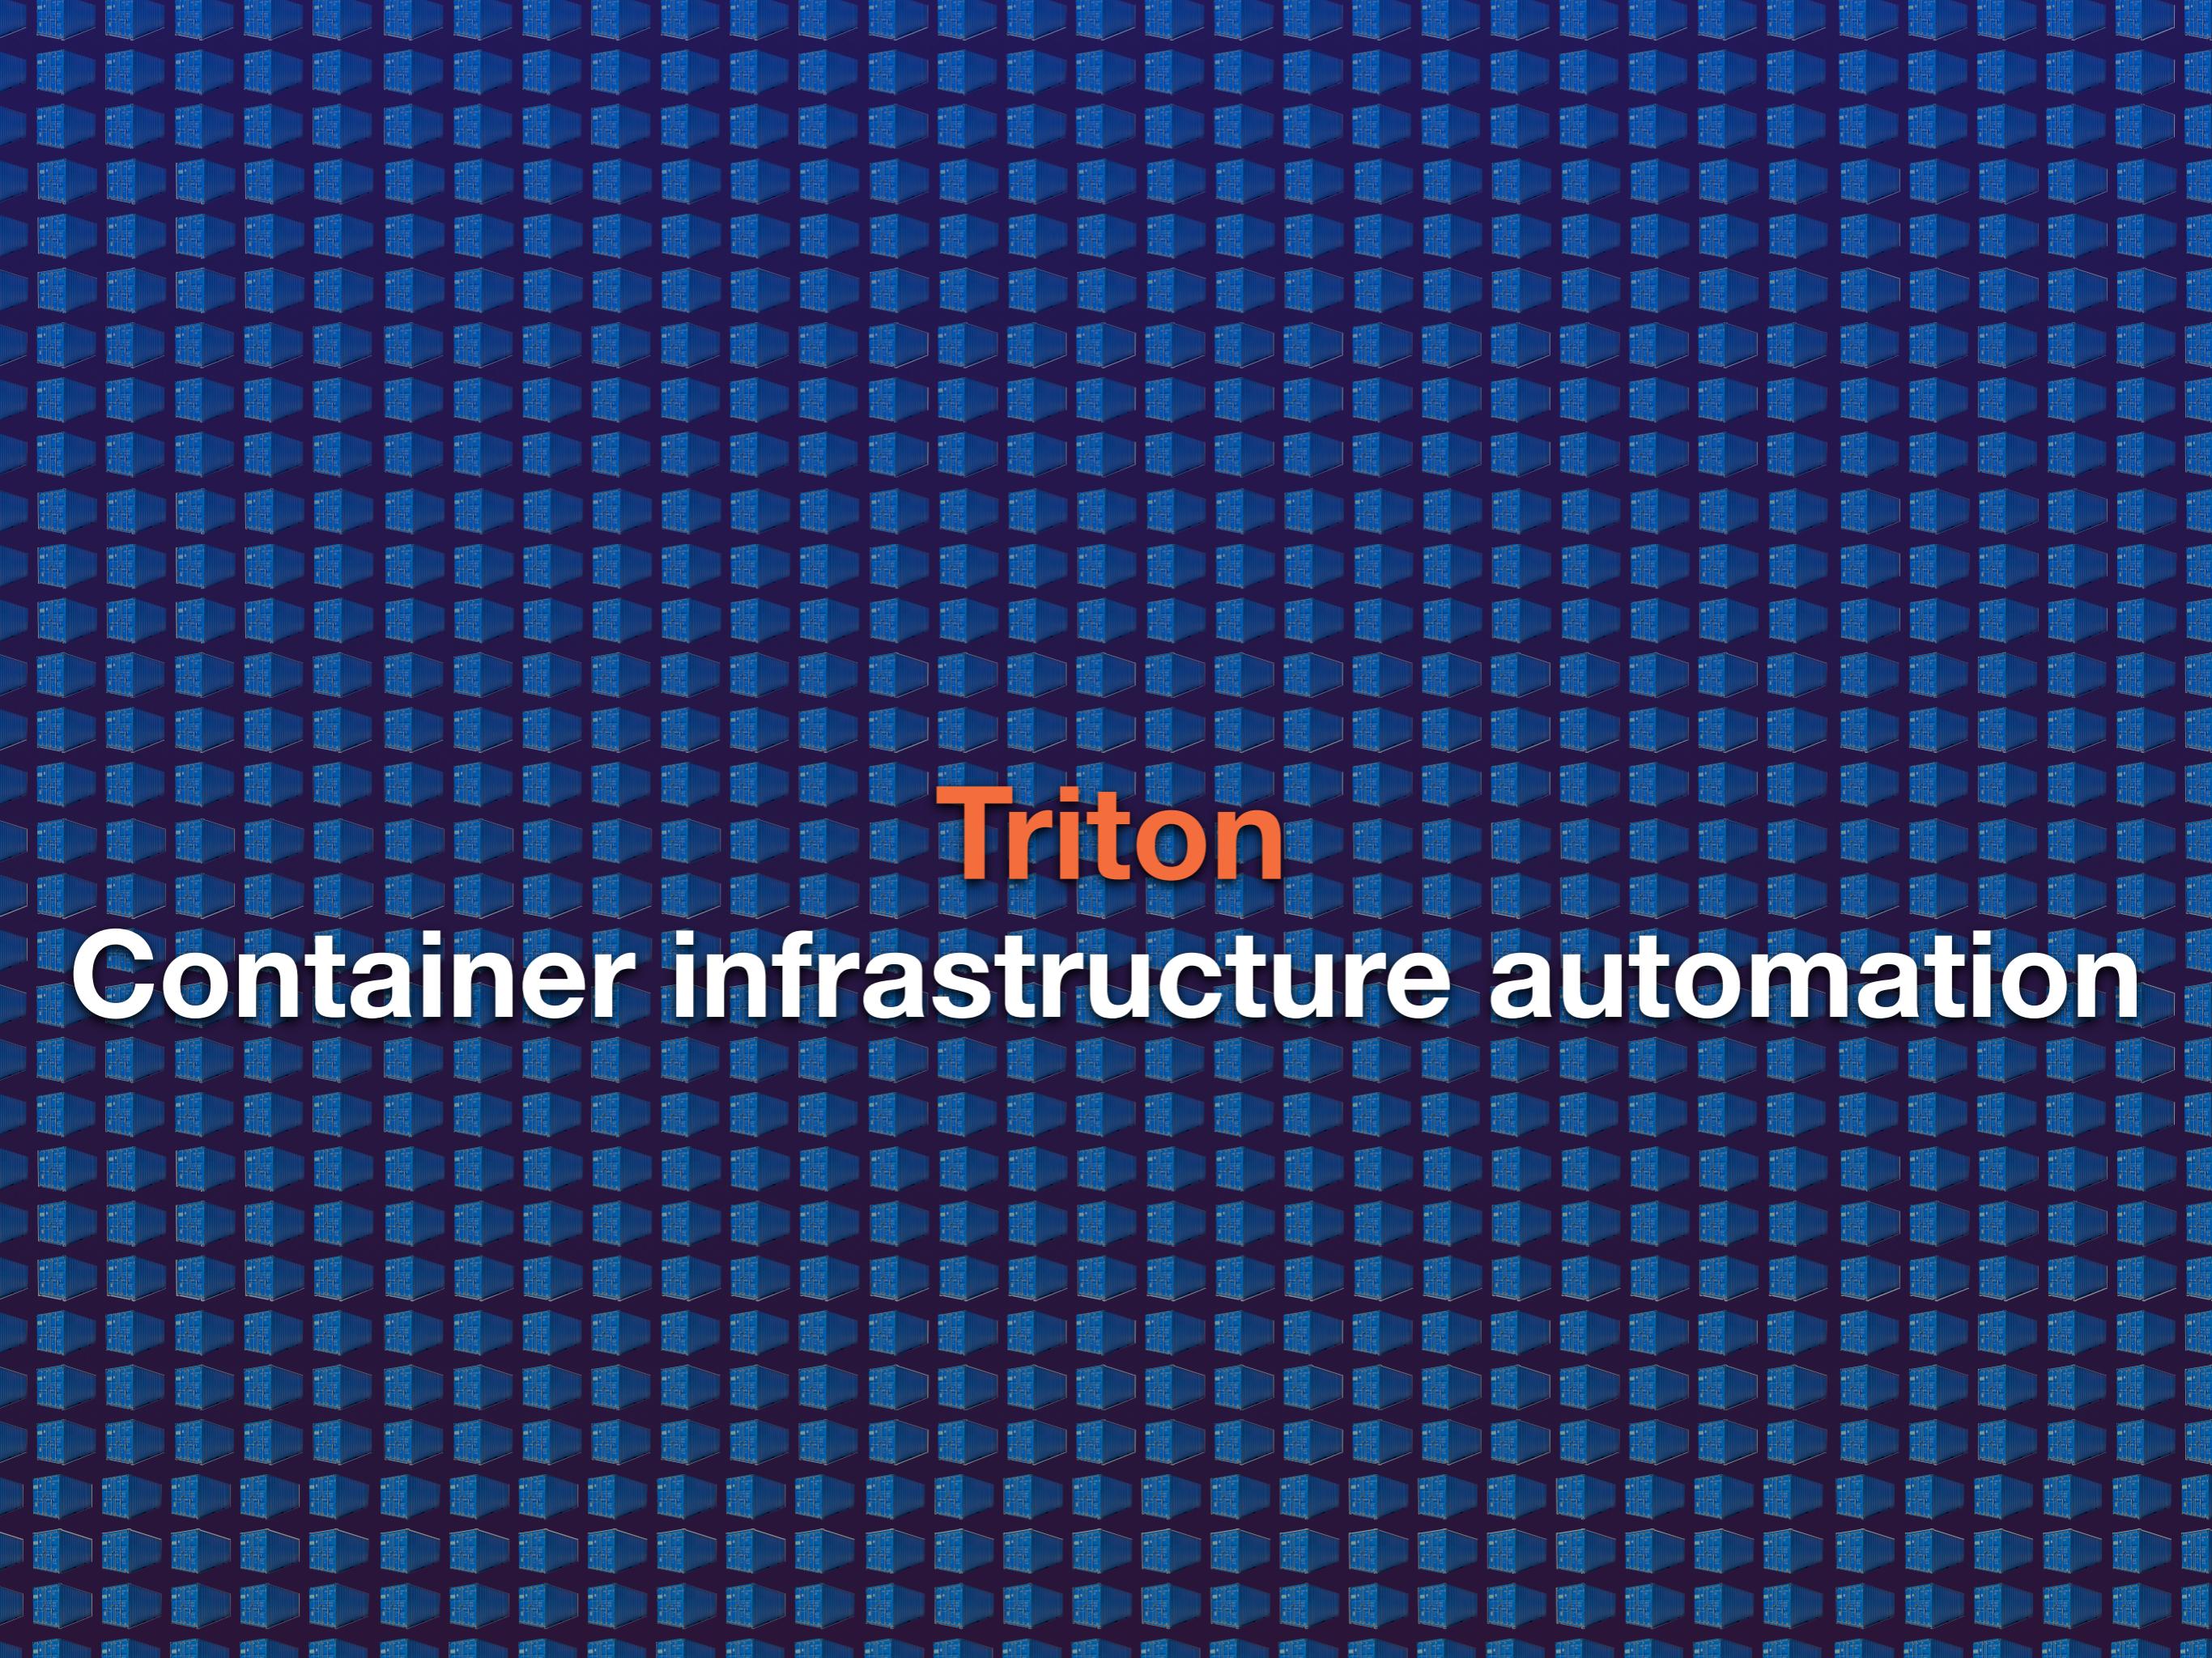
- **Open source**, Mozilla licensed
- Complete **data center automation**; including...
- Automated boot and configuration of **compute nodes**
- Multi-tenant **management layer** for compute, network, and storage across N compute nodes
- Turns an entire data center into a single, **massive container host**
- Customer-facing services, including Docker and CloudAPI to support **self-service provisioning** and management
- Bring your own application scheduling and orchestration

Infrastructure containers
Like bare metal virtual machines



SmartOS

Secure container hypervisor



Triton

Container infrastructure automation

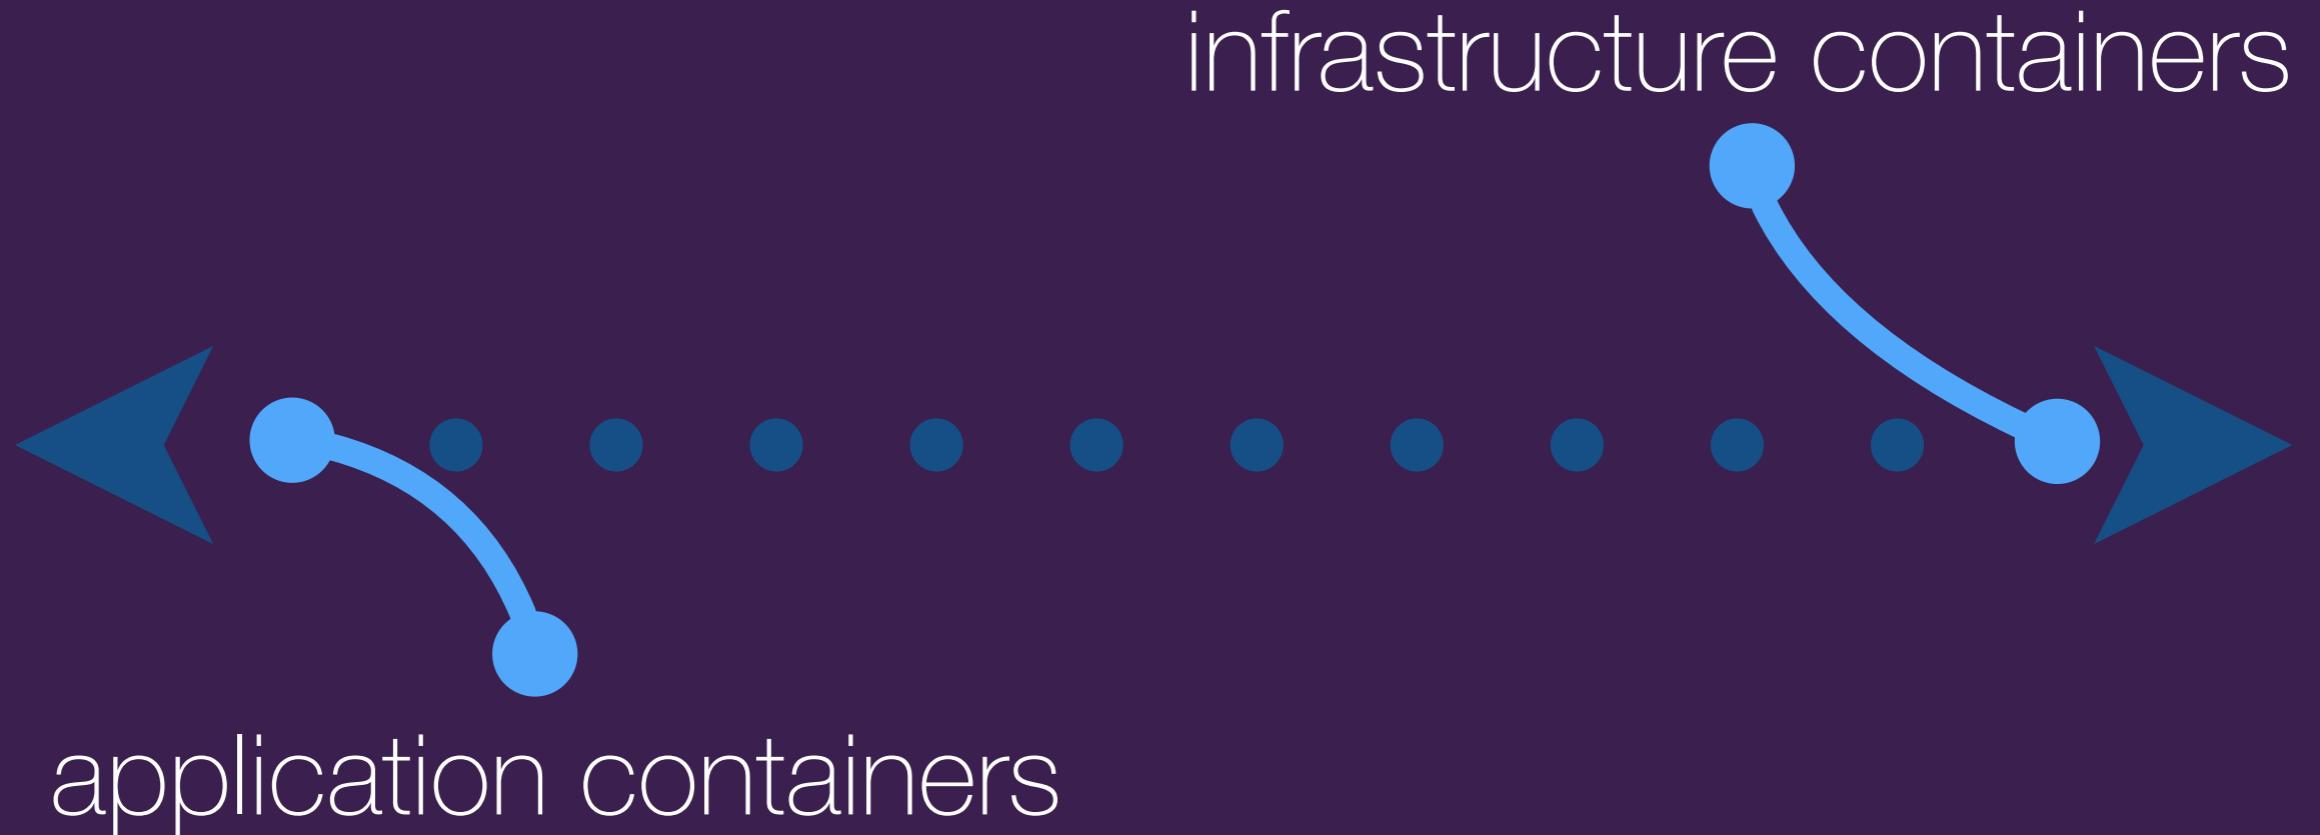
but . . .



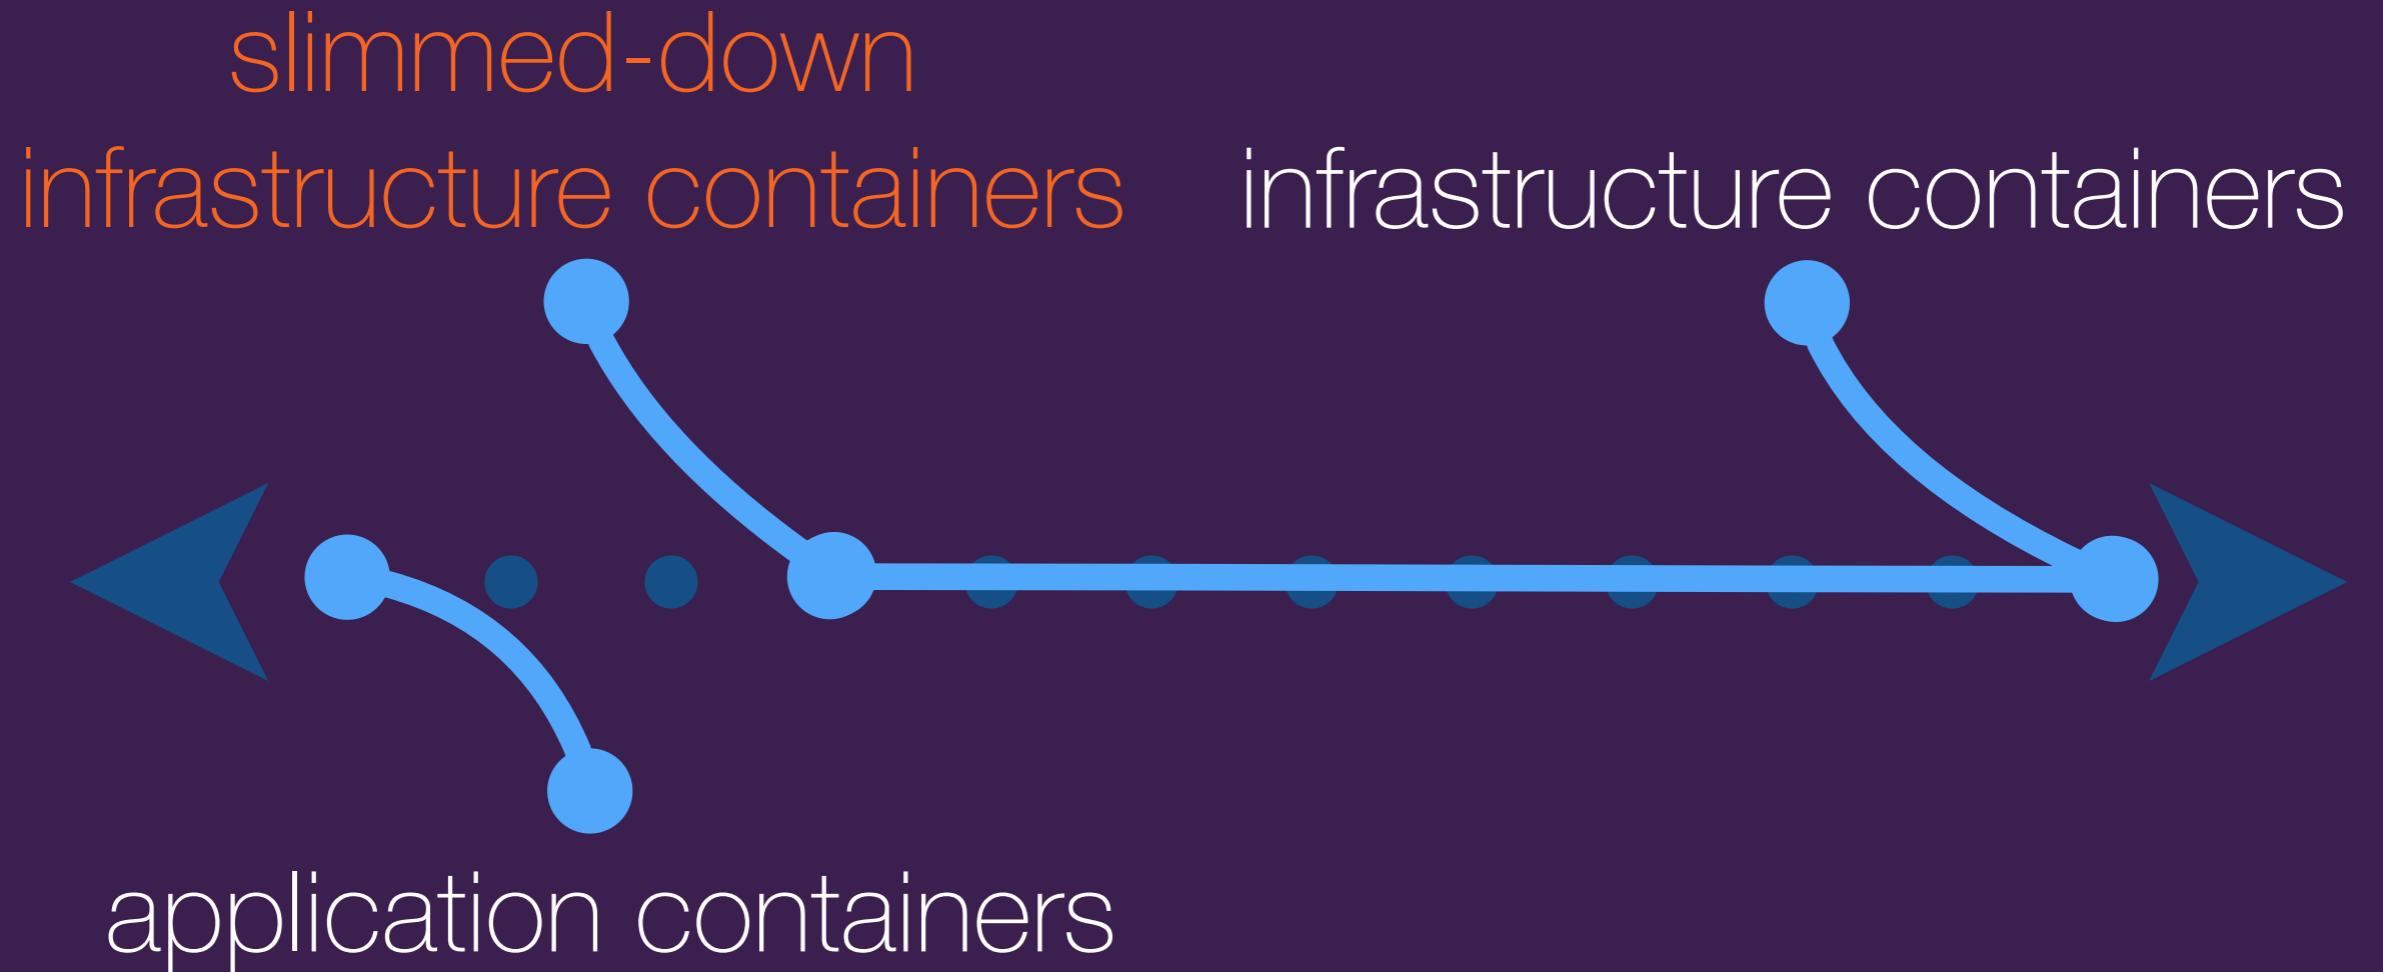
and
application containers

are
different

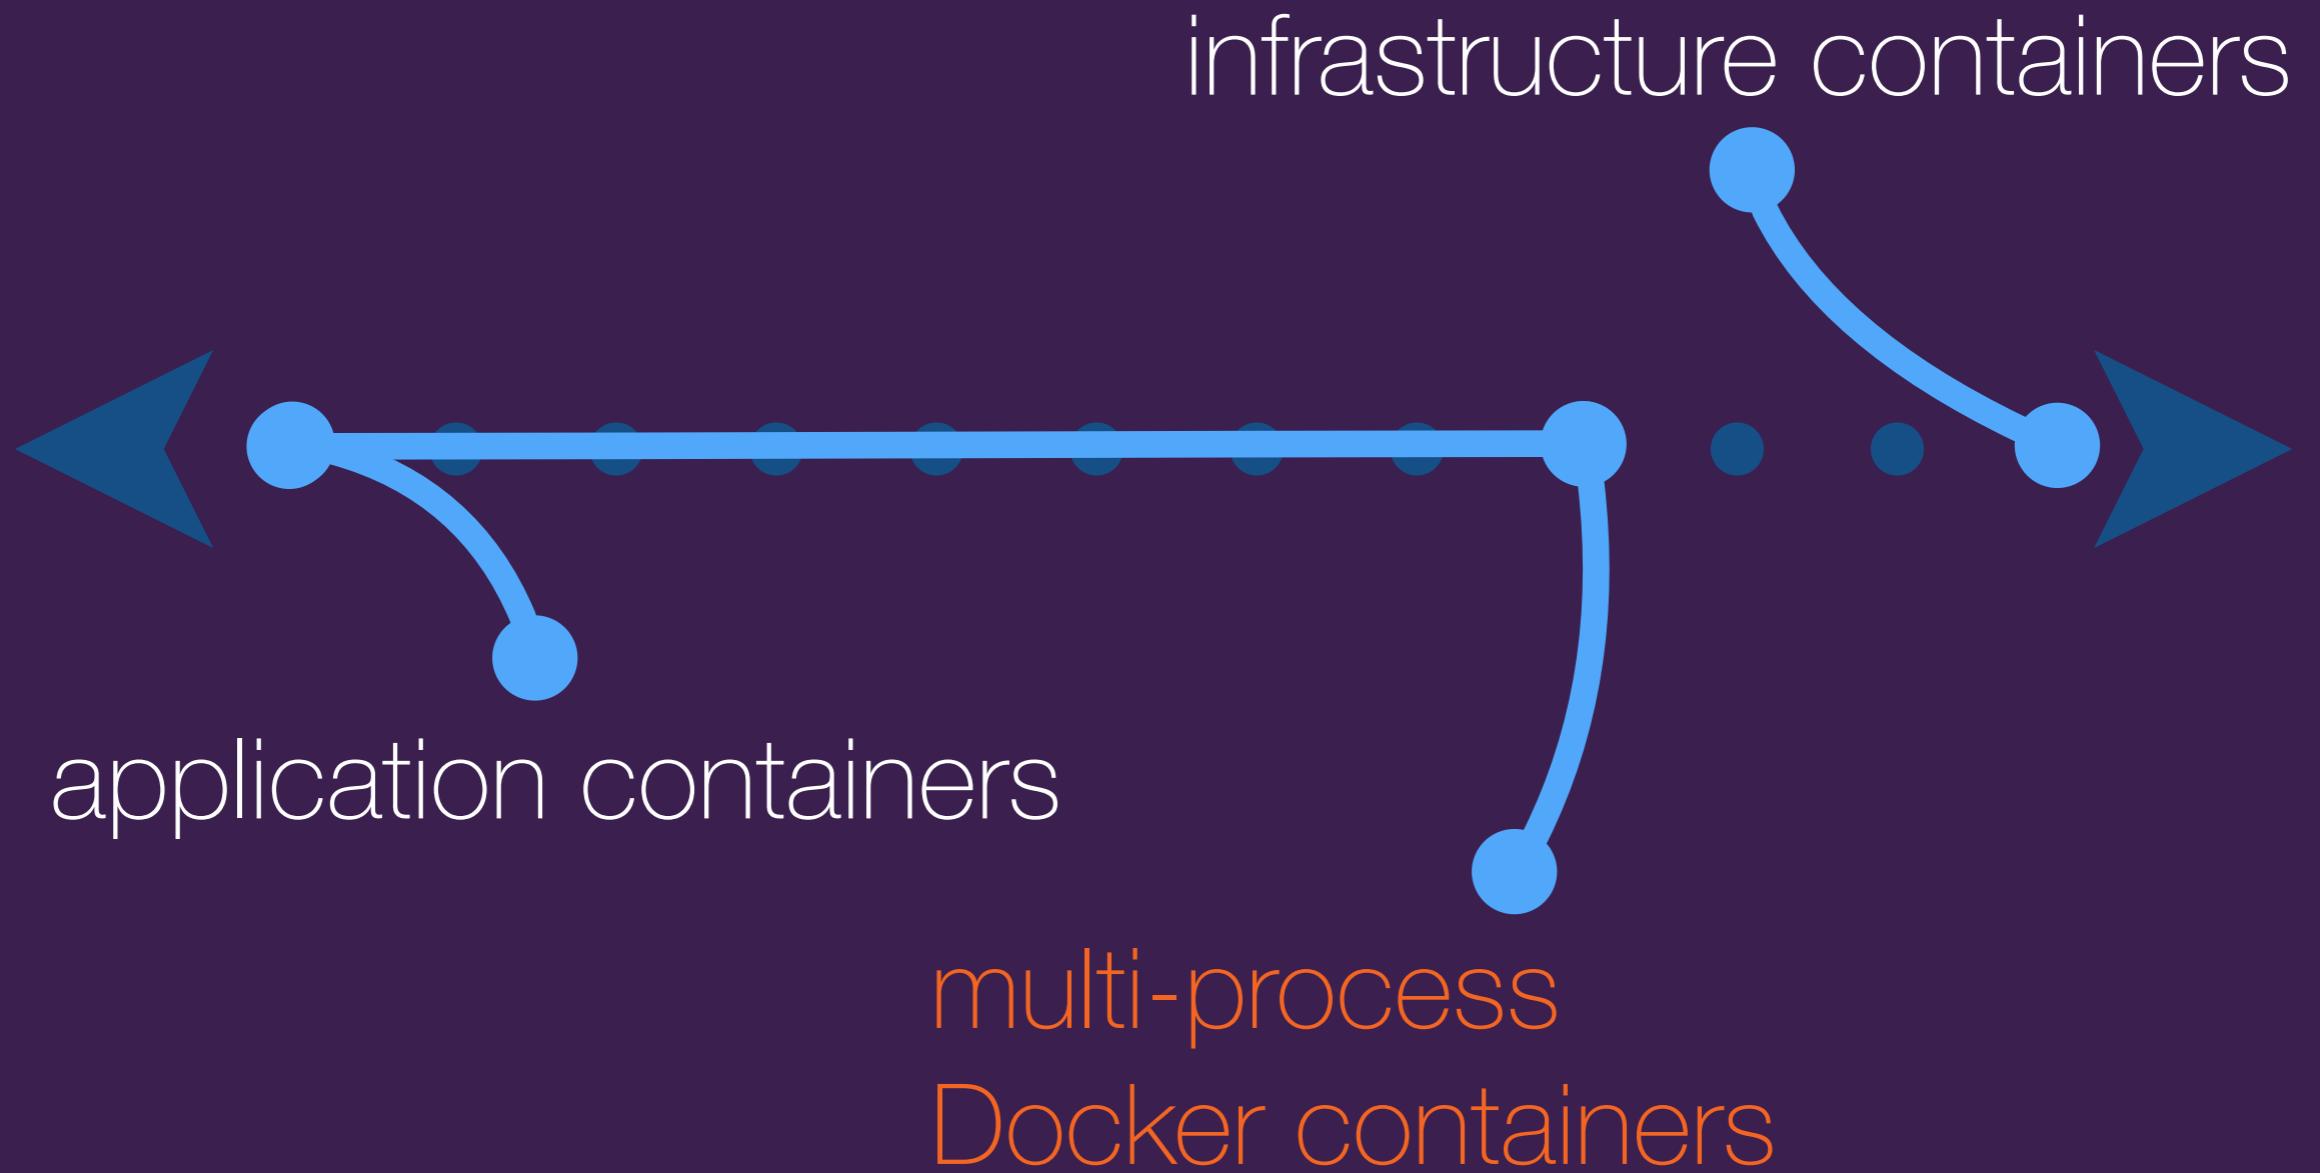
Container spectrum



Container spectrum



Container spectrum



application containers
are an opinionated
implementation

Docker container advantages

- Convenient imaging toolchain
- Easy sharing and re-use of images
- Dockerfiles are **machine+human readable** documentation for everything needed to build and run the application
- Docker Compose yaml's are machine+human readable documentation for **how to assemble a set of containers** into a larger application
- Docker begs us to **automate discovery** and configuration
- Docker allow us to **imagine deployment and execution as a single step**

Docker commoditizes the possibility of

- immutable infrastructure
- automated, testable, repeatable deploys
- the kind of apps we all dream of

Featured Post

StackShare helps you find companies using the best tools



[Sign up / Login](#)

How 500px serves up over 500TB of high res photos



500px

General Architecture

The architecture of 500px can be thought of as a large [Ruby on Rails](#) monolith surrounded by a constellation of microservices. The Rails monolith serves the main 500px web application and the 500px API, which powers the web app, mobile apps and all our third party API users. The various microservices provide specific functionality of the platform to the monolith, and also serve some API endpoints directly.

The Rails monolith is a fairly typical stack: the App and API servers serve requests with [Unicorn](#), fronted by [Nginx](#). We have clusters of these Rails servers running behind either [HAProxy](#) or LVS load balancers, and the monolith's primary datastores are [MySQL](#), [MongoDB](#), [Redis](#) and [Memcached](#). We also have a bunch of [Sidekiq](#) servers for background task processing. We currently host all of this on bare metal servers in a datacenter.

The microservices are a bit more interesting. We have about ten of them at the moment, each centered around providing an isolated and distinct business capability to the platform. Some of our microservices are:

- Search related services, built on [Elasticsearch](#)
- Content ingestion services, in front of [S3](#)
- User feeds and activity streams, built on [Roshi](#) and [AWS Kinesis](#)
- A dynamic image resizing and watermarking service
- Specialized API frontends for our web and mobile applications

We run our microservices in [Amazon EC2](#) and in our datacenter environment. They are mostly written in [Go](#), though there are a couple outliers which use [NodeJS](#) or [Sinatra](#). However, regardless of the language in use, we try to make all of our microservices good [12-factor](#) apps,

And we're not done with this problem yet! There are almost certainly more optimizations to be found, and we hope to keep pushing those response times down further and further in the future.

Workflow

We use [Github](#) and practice continuous integration for all of our primary codebases.

For the Rails monolith, we use [Semaphore](#) and [Code Climate](#). We use a standard rspec setup for unit testing, and a smaller set of Capybara/Selenium tests for integration testing. Fellow 500pxer and professional cool guy [Devon Noel de Tilly](#) has written [at length](#) about how we use those tools, so I won't try to out do him -- just go check it out.

For our Go microservices, we use [Travis CI](#) to run tests and to create Debian packages as build artifacts. Travis uploads these packages to S3, and then another system pulls them down, signs them, and imports them into our private Apt repository. We use [FPM](#) to create packages, and [Aptly](#) to manage our repos. Lately, though, I've trying out [packagecloud.io](#) and I really like it so far, so we may be changing how we do this in the near future.

For deployments, we use a combination of tools. At the lowest level we use [Ansible](#) and [Capistrano](#) for deploys and [Chef](#) for configuration management. At a higher level, we've really embraced [chatops](#) at 500px, so we've scripted the use of those tools into our beloved and loyal [Hubot](#) friend, [BMO](#).





Anyone at 500px can easily deploy the site or a microservice with a simple chat message like `bmo deploy <this thing>`. BMO goes out, deploys the thing, and then posts a log back into the chat. It's a simple, easy mechanism that has done wonders to increase visibility and reduce complexity around deploys. We use **Slack**, which is where you interact with BMO, and it makes everything really nicely searchable. If you want to find a log or if you forget how to do something, all you have to do is search the chat. Magical.

Other Important Apps

We monitor everything with **New Relic**, **Datadog**, ELK (**Elasticsearch**, **Logstash** and **Kibana**), and good old **Nagios**. We send all our emails with **Mandrill** and **Mailchimp** and we process payments with **Stripe** and **Paypal**. To help us make decisions, we use **Amazon's Elastic MapReduce** and **Redshift**, as well as **Periscope.io**. We use **Slack**, **Asana**, **Everhour**, and **Google Apps** to keep everyone in sync. And when things go wrong, we've got **Pagerduty** and **Statuspage.io** to help us out and to communicate with our users.

The Future, Comin'?

photos
or
it didn't
happen

~~photos~~

Or

~~it didn't~~

~~happen~~

repo
or
it doesn't
work

public repo
or
it doesn't
work

public repo
with

1. Dockerfile
2. docker-compose.yml
3. documentation, etc...

but . . .





pegleg sturgeon
@philsturgeon



[Follow](#)

Meanwhile, in my hometown of Bristol





Tom Jowitt

@tjbyte



[Follow](#)

@philsturgeon This is like working with Docker in production.

RETWEETS

5

FAVORITES

5



deploying
containers in VMs
slows performance

deploying
containers in VMs
adds complexity

deploying
containers in VMs
slows scaling

deploying
containers in VMs
wastes power

deploying
containers in VMs
causes global warming

deploying
containers in VMs
defeats the purpose



source



source



source



source

VMs cause complexity by...

- Requiring pre provisioning and configuration before they can host containers
- Mixing VM and container life cycles
- Reducing surface area for deployments, making pockets of unused capacity inevitable
- Reducing performance due to hardware virtualization
- Adding layers to networking
- Allowing room for us to turn them into pets
- Requiring re-packing of containers when scaling down

but . . .

containers
on bare metal Linux
struggle with
up-front costs

containers
on bare metal Linux
struggle with
elasticity

containers
on bare metal Linux
struggle with
data center automation

containers
on bare metal Linux
struggle with
multi-tenant security

Docker has changed its security status to
It's complicated

—Docker's Jérôme Petazzoni

Bare metal Linux complexities

- No data center automation
- No security domain
- No performance isolation
- No elasticity...where do you scale to?

but . . .

there is
another way

most
infrastructure

pick
two
scenario

bare
metal

- pick two:
- elasticity
 - security
 - performance

hardware
virtual
machine

pick
two:

- elasticity
- security
- ~~performance~~

container-native
infrastructure

pick
three
scenario

bare
metal
containers

- pick three:
- elasticity
 - security
 - performance

So . . .
how do you
secure it for
bare metal?

Container anatomy



Container anatomy



Container anatomy



Container anatomy



Container anatomy



Container anatomy



Open Container Foundation

Container anatomy



So . . .

can i run my
Linux images
on Triton?

yes!

Certified Ubuntu images now optimized for Joyent Triton containers

By Canonical on 18 June 2015



- Joyent Expands Triton Elastic Container Infrastructure Beyond Docker, Adds Support for Container-Native Linux on Bare Metal
- Partners with Canonical to provide certified and supported Ubuntu images, optimized to run natively on bare metal on Triton Infrastructure Containers

ubuntu
Supported by Canonical

 Joyent®

so . . .

i
promised you
a list

1.
the unit of compute
is a container

2.

you provision
containers

3.
the containers run
on bare metal

4:
the kernel offers
real security

that means
no escape
no incursion

5.

the containers are protected
from noisy neighbors

what's the point of
48 cores of bare metal
if a single container can
dominate them all?

6:
every container gets
a vNIC

6:
every container gets
a vNIC
(or two)

6:
every container gets
a vNIC
(or more)

because a
well-connected container
is a happy container

7.

You pay for
containers

7.

You pay for
containers
(not VMs)

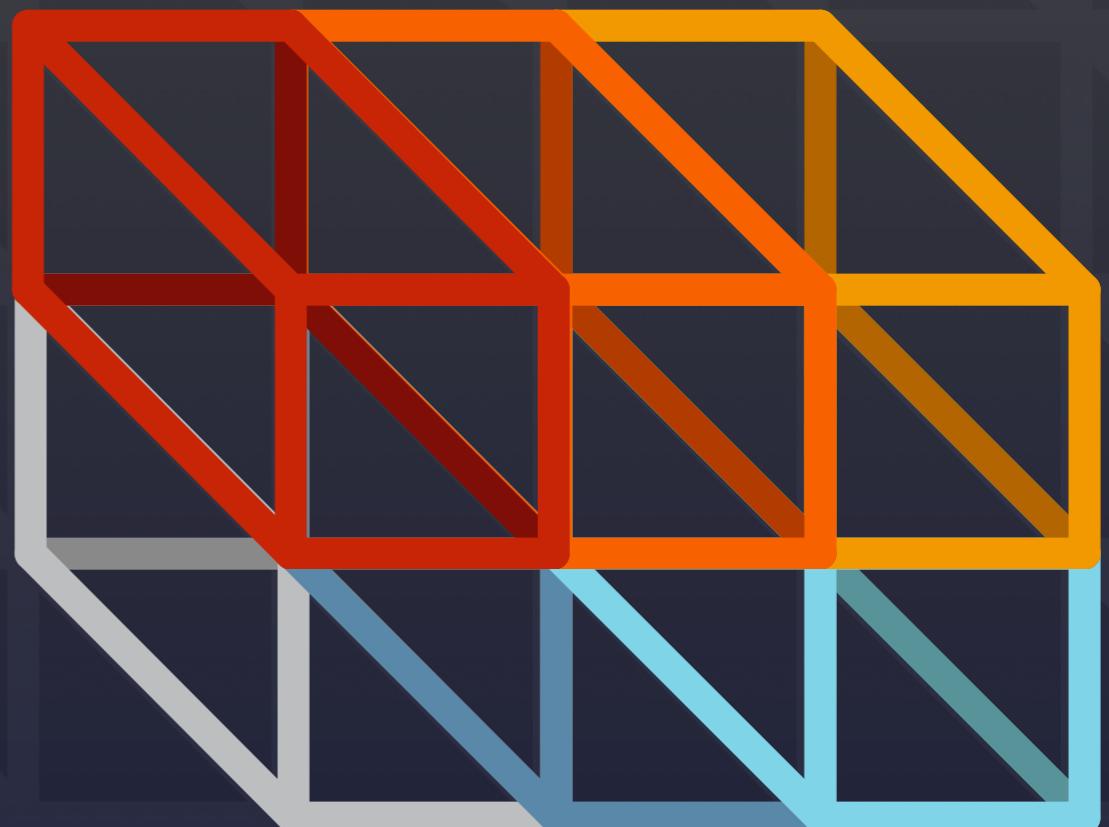
this is not
science fiction

this is
available
now



TRITON™

The best place to run containers.
Making Ops simple and scalable.



Demo
time

Container-native everything

Application architecture

- Design for automated discovery and configuration

Workflow

- Embrace containerization in every step, from dev to deploy

Infrastructure

- Automates the infrastructure from the silicon to the container
- Enforces a single lifecycle: the container
- Eliminates pet VMs and servers
- Makes provisioning and execution a single step
- Delivers the promise of containers: automatable, testable, repeatable deploys



thank you

The 7 characteristics of container-native infrastructure

1. The unit of compute is a **container**, not a VM
2. You **provision containers**, not VMs
3. The containers run on **bare metal**
4. The containers are multi-tenant bare metal **secure**
5. Every container gets its **share**
6. Every container gets one or more **VNICs**
7. You **pay for containers**, not VMs